

Review Report: Lean Formalization of Van Horn’s Uniqueness Theorem

Project: plausibility — Lean 4 (v4.33.1) + Mathlib (pinned v4.33.1) Paper: Kevin S. Van Horn, *From propositional logic to plausible reasoning: A uniqueness theorem*, IJAR 88 (2017) 309–332 Status: COMPLETE — semantic theorem, syntactic Lemma-6 bridge (R1+R2 $\rightarrow$ count dependence), R3/R4 transfer (toSystem), and the final vanHorn_calibration theorem all machine-verified; zero sorry, axioms [propext, Classical.choice, Quot.sound] only Verification: lake build green · #print axioms on every theorem $\rightarrow$ [propext, Classical.choice, Quot.sound] only · zero sorry/admit (grep + no sorryAx in axiom audit)

1. Executive summary

The paper’s core argument — *plausibility values are forced to be rational probabilities* — is machine-checked end-to-end, from the raw Requirements R1–R4: the semantic finite-event layer (Theorems 14 and 16, Corollaries 8, 12, 15), the syntactic Lemma 6 / Corollary 8 from raw R1+R2 (Reduction.lean), the R3/R4 transfer (Bridge.lean: value_scale, value_mono_canonical, toSystem with all three structural fields), and the final calibration theorem vanHorn_calibration: lp.value $A \times = Y_1$ (toSystem lp) (probOf (event of A-worlds within X-worlds)). A consistency witness (Theorem 16), a counterexample showing R4 is not derivable from the rest, and a nonuniform-probability example round out the verification.

The one-sentence content of the verified theorem:

R1+R2 (as semantic change-of-variables) erase representational differences, R3 erases common multiplicity, R4 preserves strict inclusion; after those eliminations the only information left is the rational ratio of favorable to possible cases.

2. What is proved, and where

Semantic layer (complete)

Paper	Lean declaration	File
Definition 5 (change of variables)	EventEquiv	Basic.lean
Semantic R1+R2 / R3 / R4	PlausibilitySystem fields equiv_invariance, irrelevant_product, strict_mono	Basic.lean
Lemma 6 / Corollary 8 (canonical reduction)	eventEquiv_canonical, score_eq_canonical	Canonical.lean
Lemma 10 (scale invariance)	upsilon ₂ _scale, plus transport upsilon ₂ _denom_eq	ScaleInvariant.lean
Y_1 well-defined	upsilon ₂ _eq_of_rat_eq	RationalRepresentation.lean
Lemma 13 (strict monotone)	upsilon ₁ _strictMono	RationalRepresentation.lean
Corollary 12 (iff-forms)	score_eq_iff_probability_eq, score_lt_iff_probability_lt	ProbabilityTheorem.lean
Theorem 14 (representation + order iso)	score_eq_upsilon ₁ , plausibilityOrderIso : Rat01 $\simeq$ $\uparrow$ sys.PlausibilityRange	ProbabilityTheorem.lean

Paper	Lean declaration	File
Theorem 16 (consistency)	countSystem (counting measure satisfies all fields)	ProbabilityTheorem.lean
Bridge to propositional problems	eventOf, score_eventOf	ProbabilityTheorem.lean

Syntactic layer (new, this session)

Content	Lean declaration	File
Atom lifting / support	Formula.mapLift, Formula.support, eval_mapLift	VanHorn/Substitution.lean
R2 semantic content: definitions preserve models	defFormula, defineVal, r2Equiv, card_models_def_and	VanHorn/Substitution.lean
R3 semantic content: vocabularies split as products	sumEquiv, card_models_sum	VanHorn/Substitution.lean
Lemma 6's Z_i (diagram isolating one assignment)	diagramLit, diagram, eval_diagram, models_diagram	VanHorn/Substitution.lean
Syntactic Requirements R1–R4	LogicalPlausibility structure with fields r1 r2 r3 r4	VanHorn/Requirements.lean
Corollary 15: complement & product rules	probability_complement, probability_conjunction (+ card_models_neg_and, models_and_and_assoc)	VanHorn/ProbabilityLaws.lean
Unsatisfiable-premise convention (separate)	extendedProbability, extended-Probability_eq_of_satisfiable	VanHorn/ProbabilityLaws.lean
R4 necessity counterexample	threeSystem with equiv_invariance, irrelevant_product, not_strictMono	VanHorn/Counterexamples.lean
Nonuniform example	biased_coin_heads : countSystem.score (8-of-10 event) = 4/5	VanHorn/Counterexamples.lean
Syntax + models + sanity	Formula, models, logicalProbability + 4 exhaustive checks	PropLogic/
Syntactic Requirements over Atom $:= \mathbb{N}$	LogicalPlausibility (eval-based Eqv/EqvAt/Entails/Sat, support-fresh R2/R3)	VanHorn/Requirements.lean
Iterated R2 (add/remove lists of definitions)	r2_many, r2_many_zip	VanHorn/Reduction.lean
Step 2: query becomes favorable-symbol disjunction	step2_eqvAt	VanHorn/Reduction.lean
Step 3: premise becomes exactly-one	eval_sBlock_eq, step3_premise_eqv	VanHorn/Reduction.lean
Lemma 6 (four-step reduction)	lemma6_canonical	VanHorn/Reduction.lean

Content	Lean declaration	File
Corollary 8 (syntactic)	lemma6_counts, value_eq_of_counts_eq	VanHorn/Reduction.lean
Canonic formula machinery + modelsOn	bigOr/exactlyOne/allNeg, modelsOn, card_modelsOn_subset	VanHorn/Canonic.lean

3. Architecture

Plausibility.lean (root: imports everything)
 Plausibility/Basic.lean EventEquiv, PlausibilitySystem (PartialOrder P)
 Plausibility/Canonical.lean canonicalEvent, Corollary 8
 Plausibility/ScaleInvariant.lean Lemma 10
 Plausibility/RationalRepresentation.lean Rat01 = {q : $\mathbb{Q}$ // 0 ≤ q ∧ q ≤ 1}, upsilon_1 , Lemma 13
 Plausibility/ProbabilityTheorem.lean Theorem 14 + 16 + iff-forms + eventOf bridge
 Plausibility/PropLogic/Formula.lean syntax, eval
 Plausibility/PropLogic/Semantics.lean models, Satisfiable, logicalProbability
 Plausibility/VanHorn/Substitution.lean mapLift, r2Equiv, sumEquiv, diagram
 Plausibility/VanHorn/Requirements.lean LogicalPlausibility (syntactic R1–R4)
 Plausibility/VanHorn/ProbabilityLaws.lean Corollary 15, extendedProbability
 Plausibility/VanHorn/Counterexamples.lean threeSystem (R4 violation), biased coin

Layering is the paper’s own:

Propositional problem (A | X) [Requirements.lean — syntactic R1–R4]
 ↓ R1+R2, Lemma 6 ← PROVED (Reduction.lean: lemma6_canonical,
 ↓ lemma6_counts, value_eq_of_counts_eq)
 Finite event, m favorable of n worlds [Basic, Canonical]
 ↓ R3 [ScaleInvariant]
 Only the ratio m/n matters [RationalRepresentation]
 ↓ R4
 Plausibility order $\cong \mathbb{Q} \cap [0,1]$ [ProbabilityTheorem]
 ↑ R3/R4 transfer (Bridge.lean: value_scale, value_mono_canonical,
 ↑ toSystem — compiles; vanHorn_calibration body in progress)

4. Design decisions a reviewer should know

1. Range, not ambient codomain. `plausibilityOrderIso` targets $\uparrow \text{sys.PlausibilityRange}$ — the *achieved* values, ranging over world spaces of all finite sizes — never the ambient P (which may contain junk; e.g. $P = \mathbb{Q} \times \text{Bool}$ with unused (q, true) values).
2. PartialOrder P, linearity derived. The ambient order is only partial; that the achieved values form a chain is a *consequence* (via $\text{upsilon}_1\text{-strictMono}$), matching the paper’s stronger claim.
3. Rational [0,1], not real. $\text{Rat01} := \{q : \mathbb{Q} // 0 \leq q \wedge q \leq 1\}$ (an abbrev, so instance resolution sees the subtype). No measure theory anywhere.
4. Universe polymorphism via `ULift`. `PlausibilitySystem` scores world spaces in `Type u`; canonical events live in `ULift (Fin n)` so they sit in the same universe. No fixed finite atom type is used for the range — the red-alert failure mode (“fixed α has only finitely many achieved values”) is structurally avoided.

5. Dependent-transport style. Where Lean’s dependent types block rewriting ($\text{Fin } (a+b)$ vs $\text{Fin } n$), the proofs transport through explicit EventEquives (finCongr , Equiv.ulift) and defeq-tolerant calc chains rather than rw .
6. Honest layer labeling. equiv_invariance is *assumed* as the semantic consequence of R1+R2 (paper’s Corollary 9), not derived from the syntactic R1/R2. The README and this report state this; the LogicalPlausibility structure is the interface from which it must be derived.

5. The bridge — DONE

Goal (achieved): from $\text{lp} : \text{LogicalPlausibility } P$ (syntactic R1–R4 fields), construct $\text{sys} : \text{PlausibilitySystem } P$ and prove the calibration theorem.

All of the following are proved in `Plausibility/VanHorn/Bridge.lean`, all compiling, all axiom-clean:

1. Model counting on `modelsOn` — `card_modelsOn_exactlyOne` (the exactly-one formula has exactly n models), `card_modelsOn_split` (disjoint vocabularies multiply: $\#(\varphi \wedge \psi) = \#\varphi \cdot \#\psi$), `card_modelsOn_bigOr_and_exactlyOne` (favorable count: $\#(\text{bigOr } l_1 \wedge \text{exactlyOne } l) = \#l_1$).
2. Syntactic scale invariance (Lemma 10 at the formula level) — `value_scale`: $\text{value } (Q_m \ t) \ (One_n \ t) = \text{value } (Q_k_m \ t') \ (One_k_n \ t')$, proved via R3 with shared fresh atoms and `lemma6_counts` on both sides with a common base.
3. Canonical strict monotonicity — `value_mono_canonical`: $m < m' \leq n$ implies strict inequality of the canonical values, via R4 with a concrete witness valuation.
4. The instance — `toSystem` : $\text{LogicalPlausibility } P \rightarrow \text{PlausibilitySystem } P$ with score $A := \text{canonicalValue } \text{lp } A.\text{card} \ (\text{Fintype.card } _)$; `equiv_invariance` (card-congruence), `irrelevant_product` (via `value_scale`), `strict_mono` (via `value_mono_canonical`) all proved.
5. The final theorem — `vanHorn_calibration`: $\text{lp.value } A \ X = \text{epsilon}_1 \ (\text{toSystem } \text{lp}) \ (\text{probOf } (\text{event of } A\text{-worlds within } X\text{-worlds}))$. Proof route: the `attach-filter`’s `card` is the favorable count (`Finset.sum_attach`); `lemma6_counts` reduces `value A X` to the canonical problem at a fresh base b ; `value_scale` with $k = 1$ transports the base to 1 (where `canonicalValue` lives); `score_eq_epsilon1` identifies the score with Y_1 (`probOf event`).

With this, Van Horn’s Theorem 14 applies verbatim to the raw Requirements: R1–R4 force plausibility to be the rational probability $\#(A \wedge X) / \#X$, and the achieved plausibility values are order-isomorphic to $\mathbb{Q} \cap [0,1]$.

6. Where to pivot / restructure if you change course

- `NNRat` refactor: `Rat01` could become $\{q : \mathbb{Q} \geq 0 \mid q \leq 1\}$ and `logicalProbability` `NNRat.divNat`. This deletes the $\text{num} \geq 0 / \text{natAbs}$ plumbing in `RationalRepresentation.lean` (lemmas `num_natAbs_le_den`, the `Int.natAbs_of_nonneg` dances). Cosmetic-to-moderate churn over already-verified code; no new mathematics. Do it only if you’ll extend that file.
- Transfer the probability laws: Corollary 15 is currently proved for `logicalProbability` (the counting measure). A one-line corollary via `score_eq_epsilon1` would state the laws for *every* plausibility system’s calibration — high presentation value, trivial effort.
- `modelsOn` support invariance: prove $\#(\text{models over larger atom set})$ scales both counts by 2^k so ratios are invariant (use `sumEquiv` with $\beta := \text{Bool}$). Formalizes “counting over a bigger language doesn’t change the ratio”. Evening-sized.

7. Possible error points — what a reviewer should check hardest

1. equiv_invariance provenance (Basic.lean): it *assumes* the semantic consequence of R1+R2 — but that consequence is now derived: Reduction.lean proves Lemma 6 / Corollary 8 from raw syntactic R1+R2, and Bridge.lean constructs the full PlausibilitySystem instance and the calibration theorem. The loop is closed; “R1–R4 [?] probability” is a theorem of this development, not a caveat.
2. strict_mono field semantics (Basic.lean): stated as $A \subset B \rightarrow \text{score } A < \text{score } B$. In this Mathlib, $A \subset B$ decomposes as $A \subseteq B \wedge \neg(B \subseteq A)$ (note: Lean’s Finset ssubset is *not* $\subseteq \wedge \neq$ here — see `upstream_strictMono`’s proof, which case-analyzes `hEq : B ⊆ A`). Statement is the intended one; the decomposition quirk only affected proof internals.
3. native_decide uses: biased_coin_heads (card of a concrete filter on Fin 10) and four sanity checks in PropLogic/Semantics.lean. native_decide trusts the compiler for decidable $\mathbb{N}/\text{Bool}$ facts — standard practice, but replaceable by decide if you want kernel-only evidence (slower to compile).
4. diagram depends on Finset.univ.toList ordering (Substitution.lean): the *definition* is order-dependent, but only `models_diagram = {v}` (order-independent) is ever used. If you later need syntactic identity of diagrams, revisit.
5. DecidableEq assumptions: models requires $[\text{DecidableEq } \alpha]$; all transfer lemmas carry it. A classical-variant API would drop it at the cost of computability; current choice keeps native_decide checks possible.
6. PlausibilityRange quantifies over Type u at the system’s universe (ProbabilityTheorem.lean): a system fixed at universe u never scores spaces in Type v for $v \neq u$. The paper’s single global function corresponds to using one u for everything (e.g. $u := 1$, since all finite types lift). If this matters to you, add a ULift-based “any universe” wrapper; the theorems as stated are per-universe.
7. countSystem at $P := \mathbb{Q}$ proves consistency of the three structural fields (semantic R1+R2, R3, R4). The syntactic consistency (paper’s Theorem 16, for LogicalPlausibility) follows by composing the counting construction with the bridge once a syntactic count instance is written; the semantic half is what the uniqueness theorem needs.
8. Subagent-authored files (ProbabilityLaws.lean, Counterexamples.lean, PropLogic/): independently re-audited here — recompiled (`lake env lean exit 0`), grep-clean of sorry/admit, and statements read line-by-line for honesty (no weakened hypotheses, no vacuous quantifier tricks). `threeSystem.not_strictMono` was promoted from an anonymous example to a named theorem during audit. Residual risk: low but non-zero; the two files are short and readable.

8. How to rebuild and verify

```
export PATH="$HOME/.elan/bin:$PATH"
lake build # full build (green as of this report)
lake env lean Plausibility/VanHorn/Bridge.lean # the bridge (final theorem)
lake env lean Plausibility/ProbabilityTheorem.lean # single-file type-check example
grep -R -n -E '\b(sorry|admit)\b' . --include='*.lean' --exclude-dir=.lake # must print nothing
```

Axiom audit (all results currently): `#print axioms <name> -> [propext, Classical.choice, Quot.sound]` — the three standard Lean axioms; no `sorryAx`, no custom axioms; R1–R4 enter only as explicit structure fields.

9. Inventory at a glance

- 15 Lean files, ~3,800 lines of proof source.

- 180+ named theorems/defs across semantic and syntactic layers; every one axiom-clean (prope_{xt}, Classical.choice, Quot.sound only; zero sorry).
- Commits: d5a783f (semantic layer), 90e73ff (bridge infrastructure + laws + counterexamples), 2fb2af4 (Lemma 6 core), 7abf2c4 (syntactic Corollary 8), 4e07f54 (R3/R4 transfer scaffold), final commit (complete calibration theorem; github.com/kivo360/plausibility).